

William Huynh

whuynh23@gmail.com | 206-999-3618 | linkedin.com/in/williamhuynh | williammh.github.io | Las Vegas, NV

SUMMARY

Senior engineer who ships production ML and LLM systems end to end — from self-hosted inference and prompt-driven behavioral control to the real-time data infrastructure that feeds them. Built a multi-tenant strategy-execution engine solo, evaluating user-defined logic against live market data at sub-50ms latency. Prior AI-backed products at Microsoft and enterprise SaaS at Fortune-100 scale at Smartsheet.

TECHNICAL SKILLS

ML / AI: LLM integration & deployment, self-hosted inference (Ollama, LLaMA), prompt engineering for behavioral control, agentic workflows, NumPy, Pandas, Matplotlib

Languages: Python, TypeScript, JavaScript (ES6+), C#, SQL

Backend / Data: FastAPI, asyncio, Node.js, PostgreSQL, Redis, ETL & real-time streaming pipelines, WebSockets, lock-free IPC

Cloud / Infra: AWS, Azure, GCP, Oracle Cloud (OCI ARM64), Cloudflare, Vercel, Thunder Compute, Docker

Frontend: React, Next.js, TypeScript, TailwindCSS, Shadcn

Tools: Cursor, GitHub Copilot, Claude Code, Playwright, Vitest, Git

EXPERIENCE

Software Engineer — LiveTape.ai (Independent) — *Remote* | 09/2023 – Present

- Designed and shipped a multi-tenant strategy-execution platform solo, running in production across crypto, equities, and prediction markets on Python/FastAPI and React/TypeScript.
- Built a self-hosted inference path (Ollama/LLaMA) and prompt-engineered agents that compose a trading-logic grammar — constraining LLM output to a schema validated identically on client and server, so generated strategies can never desync from execution.
- Engineered an ETL streaming pipeline evaluating a user-composed expression tree against every incoming tick at sub-50ms internal latency, versus the once-per-bar model most tools settle for.
- Architected a lock-free, zero-copy IPC layer (seqlock over a memory-mapped ring buffer) between ingestion and evaluation — readers never block writers, no mutex or message queue.
- Ran up to 50 concurrent tenant strategies in one asyncio loop with isolated per-runtime state and no DB access on the hot path.

Software Engineer II — Smartsheet — *Remote* | 06/2021 – 09/2023

- Built and maintained core features for Smartsheet Control Center, an enterprise product used by 96 of the Fortune 100, on React, TypeScript, Node.js, and Ruby.
- Contributed reusable components to the Lodestar design system, cutting UI dev time ~20% across teams; led the i18n initiative supporting 3M+ users globally.

Software Engineer / Technical Lead — Microsoft (via Mindtree / Insight Global) — *Remote* | 01/2020 – 05/2021

- Built Power Platform products including AI Builder, exposing complex ML-backed business logic through scalable APIs for external consumption.
- Owned frontend architecture modernization and a TypeScript lint migration adopted across multiple engineering teams; achieved WCAG 2.1 compliance for critical admin workflows.

EDUCATION & CERTIFICATIONS

B.S., Public Policy, University of Arizona · Azure Fundamentals (AZ-900) · AWS ML Engineer (in progress)